



The NO-U-Turn Sample (NUTS)

Presented by Per Joachims on 19.12.2019

Research Seminar in Statistics: Bayesian computation: state of the art and
recent developments

Humboldt-University of Berlin

- ▶ *Goal*: do bayesian inference based on posterior distribution
- ▶ *Problem*: posterior distribution often only known upto a normalizing constant:

$$p(\theta|x) \propto p(x|\theta)p(\theta)$$

- ▶ *Solution*: approximate posterior by **sampling** with Markov Chain Monte Carlo (MCMC)



Intro and HMC



NUTS




Ex. and Wrap-up

- ▶ The NO-U-Turn Sample (**NUTS**) by M.D. Hoffmann and A. Gelman (2011) is an extension of Hamilton Monte Carlo (**HMC**), which is a MCMC method
- ▶ HMC generates samples by setting up and simulating Hamilton dynamics
- ▶ With \mathbf{r} = momentum vector, do repeatedly:
 1. Sample \mathbf{r} from multivariate normal
 2. Evolve θ, \mathbf{r} by simulating L steps of the dynamics of the system
 3. Accept or reject sample (similar to Metropolis)



HMC: The Leapfrog Update



```

function Leapfrog( $\theta, r, \epsilon$ )
  Set  $\tilde{r} \leftarrow r + (\epsilon/2)\nabla_{\theta}\mathcal{L}(\theta)$ .
  Set  $\tilde{\theta} \leftarrow \theta + \epsilon\tilde{r}$ .
  Set  $\tilde{r} \leftarrow \tilde{r} + (\epsilon/2)\nabla_{\theta}\mathcal{L}(\tilde{\theta})$ .
  return  $\tilde{\theta}, \tilde{r}$ .

```

► *PRO:*


- **no** random walk behaviour which makes it (1) more **efficient** and (2) allows for better **dimension-scaling** compared to Metropolis Hastings and Gibbs sampling

► *CON:*

- HMC needs to compute **gradients** -> not possible for discrete variables
- Steps size ϵ and #steps L must be **tuned** well



HMC vs. NUTS



Modified version of: <https://twitter.com/sebvettelnews/status/589693881509253160>



Intro and HMC



NUTS in a Nutshell



- ▶ Eliminates the need to hand-tune L (and ϵ) and makes it **available** to more people
- ▶ Efficiency (NUTS) \geq Efficiency (well tuned HMC)
- ▶ *Core Idea*: stop when the trajectory (path of gradient steps) starts to **turn back**
- ▶ *Why?* we do not want to “ruin” our progress of exploring the sample space made so far



Source:
<http://www.fivestaryork.com/wp-content/uploads/2014/04/Nuts01.jpg>



Definition of a U-Turn


- ▶ *Idea*: stop if the sampling iteration if the trajectory starts to turn back
- ▶ *Measure*: **Dot product!**

$$\frac{d}{dt} \frac{(\tilde{\theta} - \theta) \cdot (\tilde{\theta} - \theta)}{2} = (\tilde{\theta} - \theta) \cdot \frac{d}{dt}(\tilde{\theta} - \theta) = (\tilde{\theta} - \theta) \cdot \tilde{r}$$

- ▶ Sample backwards and forwards in time to fulfill reversibility condition



A Projectory Example



Legend

- Positions on trajectory
- Starting position
- Excluded positions because of slicing
- ⊗ Excluded positions to satisfy detailed balance
- Momentum vector
- Vector between states of subtree




► Sample **iteration**:

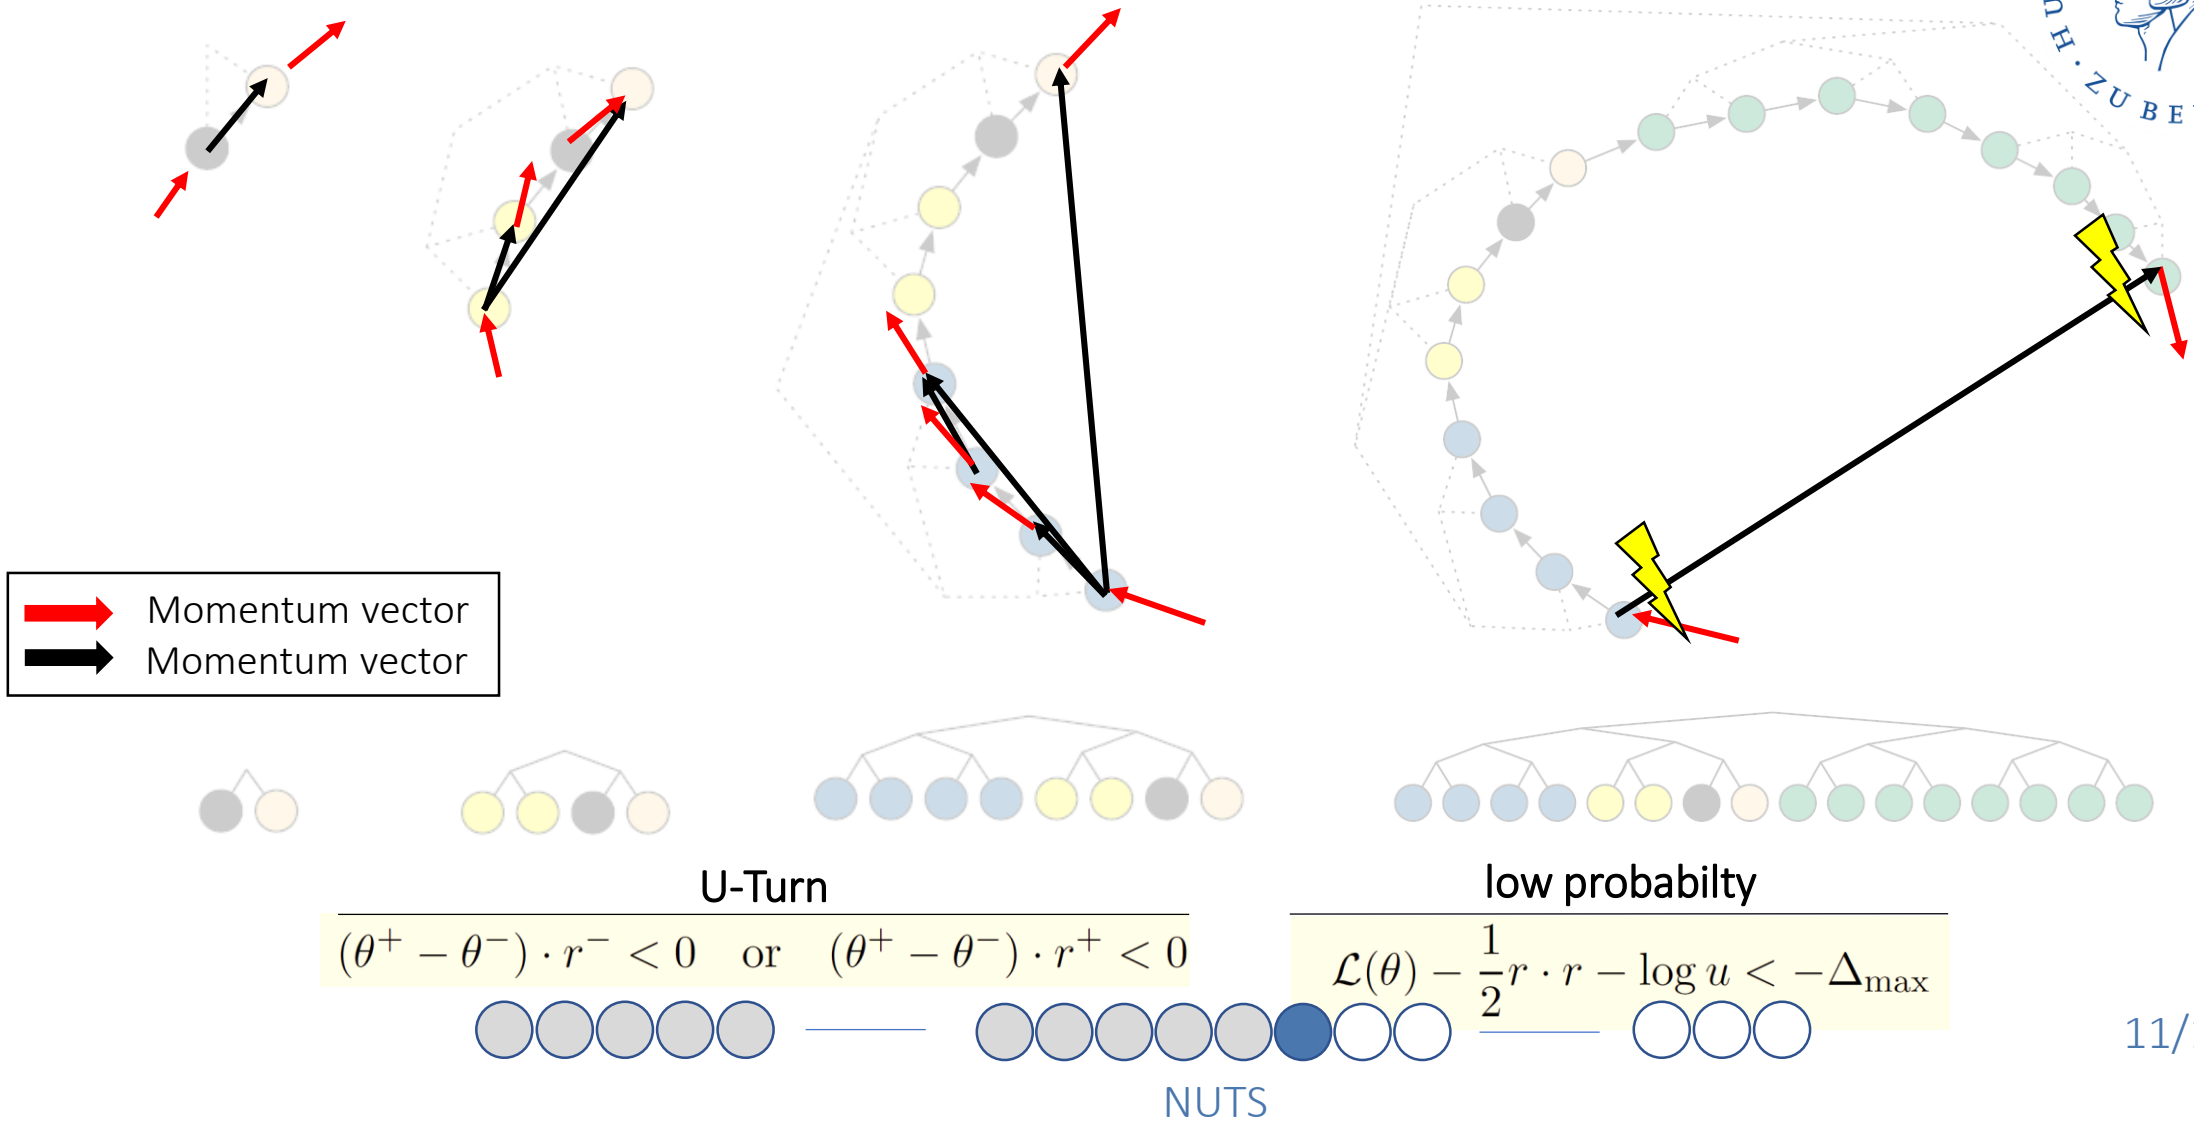
1. Sample momentum vector r
2. Sample slicing variable u
3. **Build trajectory**: trace out dynamics of θ , r forwards and backwards in time until stopping criteria (u-turn or/and probability checker) is matched
4. **Sample** uniformly from the points on trajectory



The Doubling Process



A NUTS Iteration: Applying the Stopping Criteria



<https://chi-feng.github.io/mcmc-demo/app.html>



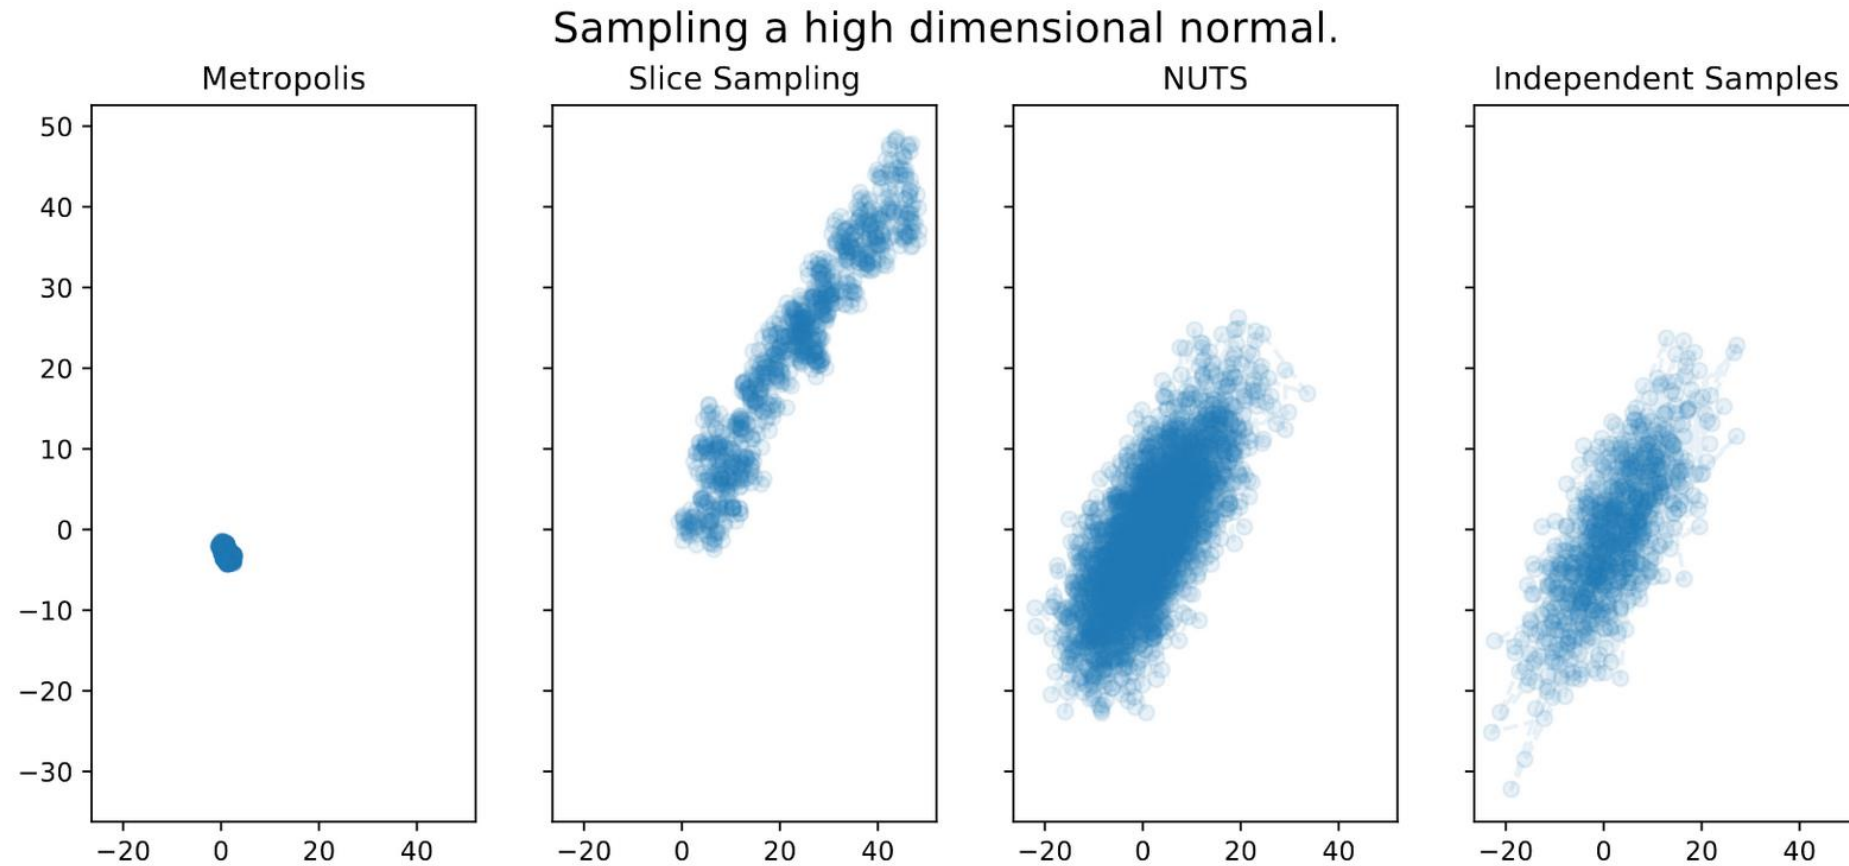
The Efficient NUTS



- ▶ (1) Stop building the trajectory if **one** leaf of the tree matches low-probability-criterion -> save *computation*
- ▶ (2) Try to sample at the end of the trajectory -> more *efficient* exploration of the state space
- ▶ (3) **sample θ , r while building** the tree to reduce size of *memory* used: $O(j) < O(2^j)$
- ▶ (4) Set step size ϵ with dual averaging



NUTS: Sampling a mvn with D=50

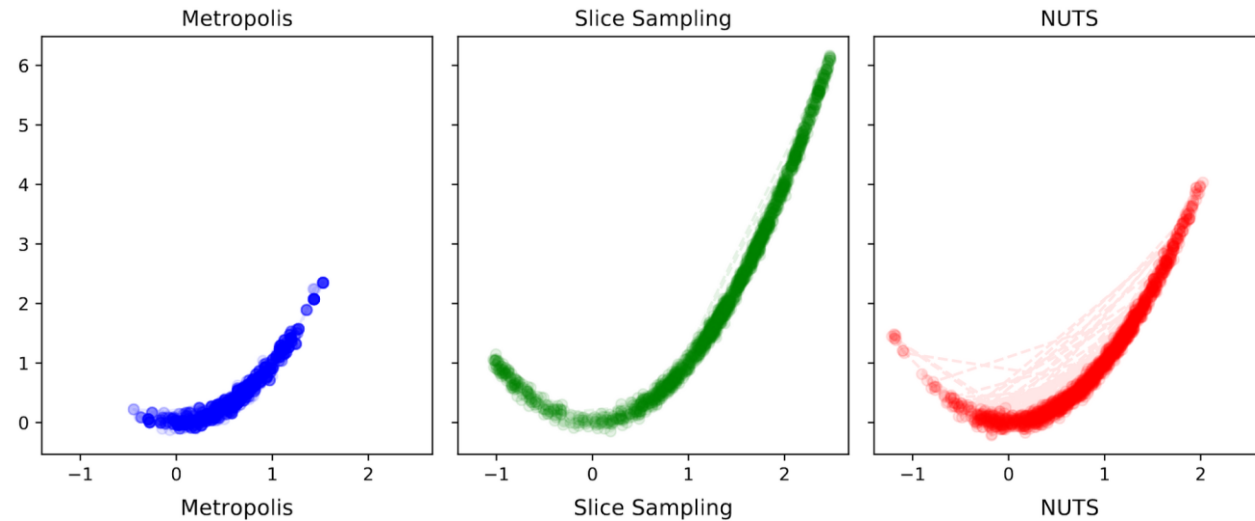



Ex. and Wrap-up

NUTS: Perks and Limitations



Rosenbrock's banana function

$$f(x, y) = (a - x)^2 + b(y - x^2)^2$$



Himmelblau's function

$$f(x, y) = (x^2 + y - 11)^2 + (x + y^2 - 7)^2$$



Ex. and Wrap-up

► *PRO:*

- no random walk behaviour which makes it (1) more **efficient** and (2) allows for better **dimension-scaling** compared to Metropolis Hastings and Gibbs sampling
- no need for handtuning

► *CON:*

- as HMC, NUTS needs to compute **gradients** -> not possible for discrete variables

➡ NUTS is default sampling method in STAN / pymc3



Ex. and Wrap-up



THANKS !

See this + code <https://github.com/pjoachims/rssbayes>


References / Further Readings



- [1] Hoffman, M.D. and Gelman, A., 2014. The No-U-Turn sampler: adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15(1), pp.1593-1623.
- [2] Betancourt, M., 2017. A conceptual introduction to Hamiltonian Monte Carlo. *arXiv preprint arXiv:1701.02434*.
- [3] Neal, R.M., 1994. An improved acceptance procedure for the hybrid Monte Carlo algorithm. *Journal of Computational Physics*, 111(1), pp.194-203.

- [4] <https://chi-feng.github.io/mcmc-demo/>
- [5] <https://docs.pymc.io/>
- [6] https://en.wikipedia.org/wiki/Test_functions_for_optimization

Appendix: HMC



Algorithm 1 Hamiltonian Monte Carlo

Given θ^0 , ϵ , L , \mathcal{L} , M :

for $m = 1$ to M **do**

 Sample $r^0 \sim \mathcal{N}(0, I)$.

 Set $\theta^m \leftarrow \theta^{m-1}$, $\tilde{\theta} \leftarrow \theta^{m-1}$, $\tilde{r} \leftarrow r^0$.

for $i = 1$ to L **do**

 Set $\tilde{\theta}, \tilde{r} \leftarrow \text{Leapfrog}(\tilde{\theta}, \tilde{r}, \epsilon)$.

end for

 With probability $\alpha = \min \left\{ 1, \frac{\exp\{\mathcal{L}(\tilde{\theta}) - \frac{1}{2}\tilde{r} \cdot \tilde{r}\}}{\exp\{\mathcal{L}(\theta^{m-1}) - \frac{1}{2}r^0 \cdot r^0\}} \right\}$, set $\theta^m \leftarrow \tilde{\theta}$, $r^m \leftarrow -\tilde{r}$.

end for

function Leapfrog(θ, r, ϵ)

 Set $\tilde{r} \leftarrow r + (\epsilon/2)\nabla_{\theta}\mathcal{L}(\theta)$.

 Set $\tilde{\theta} \leftarrow \theta + \epsilon\tilde{r}$.

 Set $\tilde{r} \leftarrow \tilde{r} + (\epsilon/2)\nabla_{\theta}\mathcal{L}(\tilde{\theta})$.

return $\tilde{\theta}, \tilde{r}$.

Appendix: Algorithm -> Naive



Algorithm 2 Naive No-U-Turn Sampler

Given $\theta^0, \epsilon, \mathcal{L}, M$:
for $m = 1$ to M do
 Resample $r^0 \sim \mathcal{N}(0, I)$.
 Resample $u \sim \text{Uniform}([0, \exp\{\mathcal{L}(\theta^{m-1} - \frac{1}{2}r^0 \cdot r^0)\}])$
 Initialize $\theta^- = \theta^{m-1}, \theta^+ = \theta^{m-1}, r^- = r^0, r^+ = r^0, j = 0, \mathcal{C} = \{(\theta^{m-1}, r^0)\}, s = 1$.
 while $s = 1$ do
 Choose a direction $v_j \sim \text{Uniform}(\{-1, 1\})$.
 if $v_j = -1$ then
 $\theta^-, r^-, -, -, \mathcal{C}', s' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v_j, j, \epsilon)$.
 else
 $-, -, \theta^+, r^+, \mathcal{C}', s' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v_j, j, \epsilon)$.
 end if
 if $s' = 1$ then
 $\mathcal{C} \leftarrow \mathcal{C} \cup \mathcal{C}'$.
 end if
 $s \leftarrow s' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$.
 $j \leftarrow j + 1$.
 end while
 Sample θ^m, r uniformly at random from \mathcal{C} .
end for

function BuildTree($\theta, r, u, v, j, \epsilon$)
if $j = 0$ then
 Base case—take one leapfrog step in the direction v .
 $\theta', r' \leftarrow \text{Leapfrog}(\theta, r, v\epsilon)$.
 $\mathcal{C}' \leftarrow \begin{cases} \{(\theta', r')\} & \text{if } u \leq \exp\{\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r'\} \\ \emptyset & \text{else} \end{cases}$
 $s' \leftarrow \mathbb{I}[\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r' > \log u - \Delta_{\max}]$.
 return $\theta', r', \theta', r', \mathcal{C}', s'$.
else
 Recursion—build the left and right subtrees.
 $\theta^-, r^-, \theta^+, r^+, \mathcal{C}', s' \leftarrow \text{BuildTree}(\theta, r, u, v, j - 1, \epsilon)$.
 if $v = -1$ then
 $\theta^-, r^-, -, -, \mathcal{C}'', s'' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v, j - 1, \epsilon)$.
 else
 $-, -, \theta^+, r^+, \mathcal{C}'', s'' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v, j - 1, \epsilon)$.
 end if
 $s' \leftarrow s' s'' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$.
 $\mathcal{C}' \leftarrow \mathcal{C}' \cup \mathcal{C}''$.
 return $\theta^-, r^-, \theta^+, r^+, \mathcal{C}', s'$.
end if

Appendix: Algorithm -> Efficient



Algorithm 3 Efficient No-U-Turn Sampler

Given $\theta^0, \epsilon, \mathcal{L}, M$:
for $m = 1$ to M do
 Resample $r^0 \sim \mathcal{N}(0, I)$.
 Resample $u \sim \text{Uniform}([0, \exp\{\mathcal{L}(\theta^{m-1} - \frac{1}{2}r^0 \cdot r^0)\}])$
 Initialize $\theta^- = \theta^{m-1}, \theta^+ = \theta^{m-1}, r^- = r^0, r^+ = r^0, j = 0, \theta^m = \theta^{m-1}, n = 1, s = 1$.
 while $s = 1$ do
 Choose a direction $v_j \sim \text{Uniform}(\{-1, 1\})$.
 if $v_j = -1$ then
 $\theta^-, r^-, -, -, \theta', n', s' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v_j, j, \epsilon)$.
 else
 $-, -, \theta^+, r^+, \theta', n', s' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v_j, j, \epsilon)$.
 end if
 if $s' = 1$ then
 With probability $\min\{1, \frac{n'}{n}\}$, set $\theta^m \leftarrow \theta'$.
 end if
 $n \leftarrow n + n'$.
 $s \leftarrow s' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$.
 $j \leftarrow j + 1$.
 end while
end for

function BuildTree($\theta, r, u, v, j, \epsilon$)
if $j = 0$ then
 Base case—take one leapfrog step in the direction v .
 $\theta', r' \leftarrow \text{Leapfrog}(\theta, r, v\epsilon)$.
 $n' \leftarrow \mathbb{I}[u \leq \exp\{\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r'\}]$.
 $s' \leftarrow \mathbb{I}[\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r' > \log u - \Delta_{\max}]$
 return $\theta', r', \theta', r', \theta', n', s'$.
else
 Recursion—implicitly build the left and right subtrees.
 $\theta^-, r^-, \theta^+, r^+, \theta', n', s' \leftarrow \text{BuildTree}(\theta, r, u, v, j - 1, \epsilon)$.
 if $s' = 1$ then
 if $v = -1$ then
 $\theta^-, r^-, -, -, \theta'', n'', s'' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v, j - 1, \epsilon)$.
 else
 $-, -, \theta^+, r^+, \theta'', n'', s'' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v, j - 1, \epsilon)$.
 end if
 With probability $\frac{n''}{n' + n''}$, set $\theta' \leftarrow \theta''$.
 $s' \leftarrow s'' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$
 $n' \leftarrow n' + n''$
 end if
 return $\theta^-, r^-, \theta^+, r^+, \theta', n', s'$.
end if

Appendix: Satisfying Detailed-Balance $p(\mathcal{B}, \mathcal{C} | \theta, r, u, \epsilon)$

C.1: All elements of \mathcal{C} must be chosen in a way that preserves volume. That is, any deterministic transformations of θ, r used to add a state θ', r' to \mathcal{C} must have a Jacobian with unit determinant.

Why?: threat unnormalized prob. density of element as unconditional probability mass

How?: leapfrog step are volume preserving

C.2: $p((\theta, r) \in \mathcal{C} | \theta, r, u, \epsilon) = 1.$

Why?: ensure current state is a valid sample

How?: satisfied if initial state is in proposal states

C.3: $p(u \leq \exp\{\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r'\} | (\theta', r') \in \mathcal{C}) = 1.$

Why?: any state in \mathcal{C} must be in the slice defined by $u \rightarrow$ all states have equal and positive cond. prob. density

How?: satisfied due to the slicing variable

C.4: If $(\theta, r) \in \mathcal{C}$ and $(\theta', r') \in \mathcal{C}$ then for any \mathcal{B} , $p(\mathcal{B}, \mathcal{C} | \theta, r, u, \epsilon) = p(\mathcal{B}, \mathcal{C} | \theta', r', u, \epsilon).$

Why?: \mathcal{B} and \mathcal{C} must have equal prob. to be selected regardless of the state

How?: exclude from \mathcal{C} any state that could not have generated \mathcal{B}

Appendix: Transition Kernel of the Efficient NUTS



$$T(w'|w, \mathcal{C}) = \begin{cases} \frac{\mathbb{I}[w' \in \mathcal{C}^{\text{new}}]}{|\mathcal{C}^{\text{new}}|} & \text{if } |\mathcal{C}^{\text{new}}| > |\mathcal{C}^{\text{old}}|, \\ \frac{|\mathcal{C}^{\text{new}}|}{|\mathcal{C}^{\text{old}}|} \frac{\mathbb{I}[w' \in \mathcal{C}^{\text{new}}]}{|\mathcal{C}^{\text{new}}|} + \left(1 - \frac{|\mathcal{C}^{\text{new}}|}{|\mathcal{C}^{\text{old}}|}\right) \mathbb{I}[w' = w] & \text{if } |\mathcal{C}^{\text{new}}| \leq |\mathcal{C}^{\text{old}}| \end{cases}$$

Appendix: Algorithm -> Differences in Main Loop



Algorithm 2 Naive No-U-Turn Sampler

Given $\theta^0, \epsilon, \mathcal{L}, M$:
for $m = 1$ to M do
 Resample $r^0 \sim \mathcal{N}(0, I)$.
 Resample $u \sim \text{Uniform}([0, \exp\{\mathcal{L}(\theta^{m-1} - \frac{1}{2}r^0 \cdot r^0)\}])$
 Initialize $\theta^- = \theta^{m-1}, \theta^+ = \theta^{m-1}, r^- = r^0, r^+ = r^0, j = 0, \mathcal{C} = \{(\theta^{m-1}, r^0)\}, s = 1$.
 while $s = 1$ do
 Choose a direction $v_j \sim \text{Uniform}(\{-1, 1\})$.
 if $v_j = -1$ then
 $\theta^-, r^-, -, -, \mathcal{C}', s' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v_j, j, \epsilon)$.
 else
 $-, -, \theta^+, r^+, \mathcal{C}', s' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v_j, j, \epsilon)$.
 end if
 if $s' = 1$ then
 $\mathcal{C} \leftarrow \mathcal{C} \cup \mathcal{C}'$.
 end if
 $s \leftarrow s' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$.
 $j \leftarrow j + 1$.
 end while
 Sample θ^m, r uniformly at random from \mathcal{C} .
end for

Algorithm 3 Efficient No-U-Turn Sampler

Given $\theta^0, \epsilon, \mathcal{L}, M$:
for $m = 1$ to M do
 Resample $r^0 \sim \mathcal{N}(0, I)$.
 Resample $u \sim \text{Uniform}([0, \exp\{\mathcal{L}(\theta^{m-1} - \frac{1}{2}r^0 \cdot r^0)\}])$
 Initialize $\theta^- = \theta^{m-1}, \theta^+ = \theta^{m-1}, r^- = r^0, r^+ = r^0, j = 0, \theta^m = \theta^{m-1}, n = 1, s = 1$.
 while $s = 1$ do
 Choose a direction $v_j \sim \text{Uniform}(\{-1, 1\})$.
 if $v_j = -1$ then
 $\theta^-, r^-, -, -, \theta', n', s' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v_j, j, \epsilon)$.
 else
 $-, -, \theta^+, r^+, \theta', n', s' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v_j, j, \epsilon)$.
 end if
 if $s' = 1$ then
 With probability $\min\{1, \frac{n'}{n}\}$, set $\theta^m \leftarrow \theta'$.
 end if
 $n \leftarrow n + n'$.
 $s \leftarrow s' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$.
 $j \leftarrow j + 1$.
 end while
end for

Appendix: Algorithm -> Differences in Tree Growing



Naive

```
function BuildTree( $\theta, r, u, v, j, \epsilon$ )
if  $j = 0$  then
    Base case—take one leapfrog step in the direction  $v$ .
     $\theta', r' \leftarrow \text{Leapfrog}(\theta, r, v\epsilon)$ .
     $\mathcal{C}' \leftarrow \begin{cases} \{(\theta', r')\} & \text{if } u \leq \exp\{\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r'\} \\ \emptyset & \text{else} \end{cases}$ 
     $s' \leftarrow \mathbb{I}[\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r' > \log u - \Delta_{\max}]$ .
    return  $\theta', r', \theta', r', \mathcal{C}', s'$ .
else
    Recursion—build the left and right subtrees.
     $\theta^-, r^-, \theta^+, r^+, \mathcal{C}', s' \leftarrow \text{BuildTree}(\theta, r, u, v, j-1, \epsilon)$ .
    if  $v = -1$  then
         $\theta^-, r^-, -, -, \mathcal{C}'', s'' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v, j-1, \epsilon)$ .
    else
         $-, -, \theta^+, r^+, \mathcal{C}'', s'' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v, j-1, \epsilon)$ .
    end if
     $s' \leftarrow s' s'' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$ .
     $\mathcal{C}' \leftarrow \mathcal{C}' \cup \mathcal{C}''$ .
    return  $\theta^-, r^-, \theta^+, r^+, \mathcal{C}', s'$ .
end if
```

Efficient

```
function BuildTree( $\theta, r, u, v, j, \epsilon$ )
if  $j = 0$  then
    Base case—take one leapfrog step in the direction  $v$ .
     $\theta', r' \leftarrow \text{Leapfrog}(\theta, r, v\epsilon)$ .
     $n' \leftarrow \mathbb{I}[u \leq \exp\{\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r'\}]$ .
     $s' \leftarrow \mathbb{I}[\mathcal{L}(\theta') - \frac{1}{2}r' \cdot r' > \log u - \Delta_{\max}]$ 
    return  $\theta', r', \theta', r', \theta', n', s'$ .
else
    Recursion—implicitly build the left and right subtrees.
     $\theta^-, r^-, \theta^+, r^+, \theta', n', s' \leftarrow \text{BuildTree}(\theta, r, u, v, j-1, \epsilon)$ .
    if  $s' = 1$  then
        if  $v = -1$  then
             $\theta^-, r^-, -, -, \theta'', n'', s'' \leftarrow \text{BuildTree}(\theta^-, r^-, u, v, j-1, \epsilon)$ .
        else
             $-, -, \theta^+, r^+, \theta'', n'', s'' \leftarrow \text{BuildTree}(\theta^+, r^+, u, v, j-1, \epsilon)$ .
        end if
        With probability  $\frac{n''}{n' + n''}$ , set  $\theta' \leftarrow \theta''$ .
         $s' \leftarrow s'' \mathbb{I}[(\theta^+ - \theta^-) \cdot r^- \geq 0] \mathbb{I}[(\theta^+ - \theta^-) \cdot r^+ \geq 0]$ 
         $n' \leftarrow n' + n''$ 
    end if
    return  $\theta^-, r^-, \theta^+, r^+, \theta', n', s'$ .
end if
```